

## Session 12 – Oozie and Flume Assignment1

### Streaming Twitter Data Using Flume

#### Streaming Twitter Data

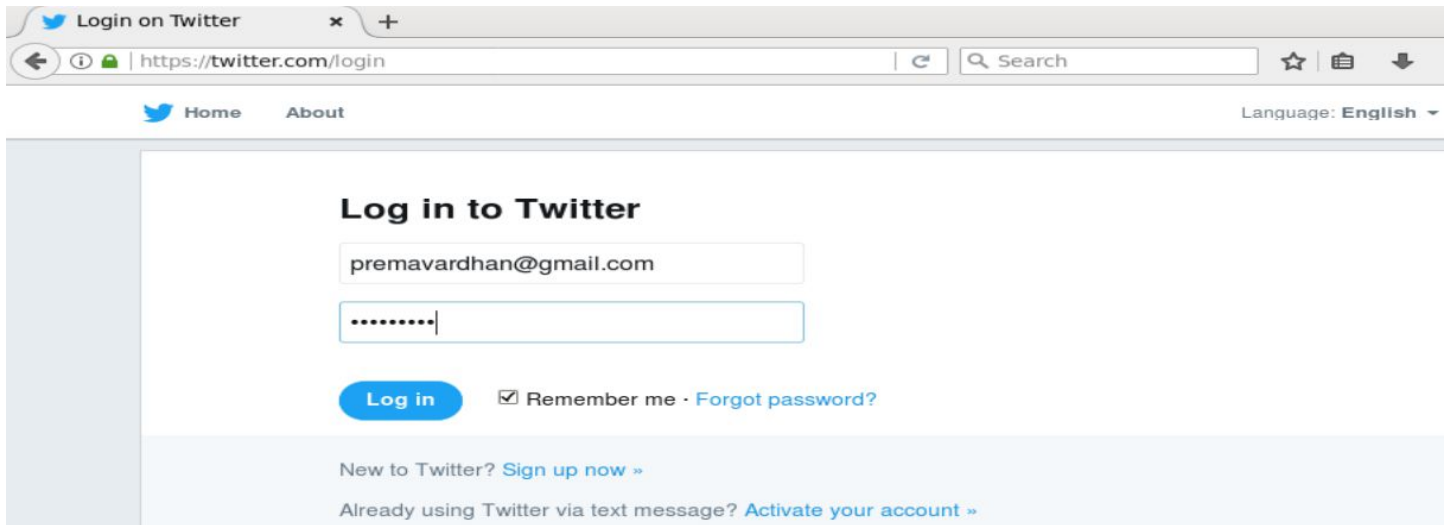
To stream data to our database from twitter we should have the following pre-requisites.

- Twitter account.
- Hadoop cluster.

If both prerequisites are available we can move to our further step.

#### Step 1:

Logging to the twitter account



#### Step 2:

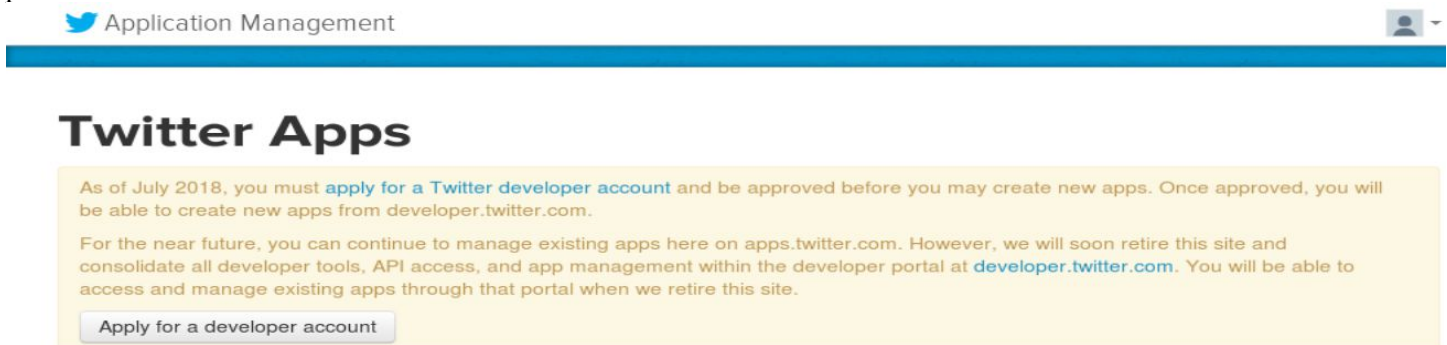
Go to the following link and click the 'create new app' button.

<https://apps.twitter.com/app>


You  
picture.

get

following



Click on the 'Apply for a development account.


Developer

Use cases


Products


Docs


More


Apply


STATUS: IN PROGRESS

User profile

Account details

Use case details

Terms of service

Email verification


Interested in a developer account?

Some of our premium APIs are currently in Beta. By applying, you agree to receive emails from our team requesting feedback on your experience.

Select a user profile to associate


By default, this @username will be the admin of this developer account. If you are creating a developer account on behalf of your organization, you may wish to use your organization's @username as it is most likely to own the Apps you will use to access the API endpoints or warrant special permissions. You'll be able to invite teammates and re-assign roles later within your developer account settings.

Associate your current Twitter @username

premavardhanreddy  
@premavardhan

Continue

Now click on continue button.


Developer


Use cases

Products

Docs

More

DashboardPrema Vardhan




Create your first app

You'll need an app and API key in order to authenticate and integrate with most Twitter developer products. Create an app to get your API key.

Create an app

**Step 3:**  
Enter the necessary details.

Developer

Use cases


Products


Docs

More

Dashboard

Prema Vardhan



Understanding apps

What is an app?

Why register an app?

Which products require an API key?

### App details

The following app details will be visible to app users and are required to generate the API keys needed to authenticate Twitter developer products.

App name (required)

Acadgild-PremaApp

Maximum characters: 32

Application description (required)

Share a description of your app. This description will be visible to users so this is a good place to tell them what your app does.


This App will help me do analysis in flume.

Between 10 and 200 characters

Website URL (required)

https://yahoo.com

**Step 4:**  
Select the 'keys and Access Token' tab.

Developer

Use cases


Products

Docs

More

Dashboard

Prema Vardhan



Apps / Acadgild-PremaApp

Create an app

App details

Keys and tokens

Permissions

### Keys and tokens

Keys, secret keys and access tokens management.

Consumer API keys

TE3NZCwVDMqoRPfs8RxADyiHb (API key)

BTYeE2rVmS96EWAq7KY3OOywpIO6O1swUsAoFQd7lcsQwU7HG3 (API secret key)

Regenerate

Access token & access token secret

None

Create

**Step 5:**  
Copy the consumer key and the consumer secret key:

**Step 6:**  
Copy the Flume configuration code from the below link and pass it in the newly created file:  
<https://drive.google.com/open?id=0B1QaXx7tpw3Sb3U4LW9SWINidkk>

```

TwitterAgent.sources = Twitter
TwitterAgent.channels = MemChannel
TwitterAgent.sinks = HDFS

# Describing/Configuring the source
TwitterAgent.sources.Twitter.type = org.apache.flume.source.twitter.TwitterSource
TwitterAgent.sources.Twitter.consumerKey=uX0TWqkx0okYEjjqLzxIx6mD6
TwitterAgent.sources.Twitter.consumerSecret=rzHIs3TMJnADbZNvdGU7LQUo0kPxPISq3RGSLfqcBip39X5END
TwitterAgent.sources.Twitter.accessToken=559516596-yDA9xq0ljo4CV32wSnqxs2BXh4RBIRKFxZG5ZrPC
TwitterAgent.sources.Twitter.accessTokenSecret=zDxePILZitS5tIWbhre0GWqps0FIj90adX8RZb6w8ZCwz
TwitterAgent.sources.Twitter.keywords=hadoop, bigdata, mapreduce, mahout, hbase, nosql
# Describing/Configuring the sink

TwitterAgent.sources.Twitter.keywords= hadoop,election,sports, cricket,Big data

TwitterAgent.sinks.HDFS.channel=MemChannel
TwitterAgent.sinks.HDFS.type=hdfs
TwitterAgent.sinks.HDFS.hdfs.path=hdfs://localhost:9000/user/flume/tweets
TwitterAgent.sinks.HDFS.hdfs.fileType=DataStream
TwitterAgent.sinks.HDFS.hdfs.writeformat=Text
TwitterAgent.sinks.HDFS.hdfs.batchSize=1000
TwitterAgent.sinks.HDFS.hdfs.rollSize=0
TwitterAgent.sinks.HDFS.hdfs.rollCount=10000
TwitterAgent.sinks.HDFS.hdfs.rollInterval=600

TwitterAgent.channels.MemChannel.type=memory
TwitterAgent.channels.MemChannel.capacity=10000
TwitterAgent.channels.MemChannel.transactionCapacity=1000

TwitterAgent.sources.Twitter.channels = MemChannel
TwitterAgent.sinks.HDFS.channel = MemChannel

```

### Step 7:

Change the twitter api keys with the keys generated as we were created the twitter app.

```

TwitterAgent.sources = Twitter
TwitterAgent.channels = MemChannel
TwitterAgent.sinks = HDFS

# Describing/Configuring the source
TwitterAgent.sources.Twitter.type = org.apache.flume.source.twitter.TwitterSource
TwitterAgent.sources.Twitter.consumerKey=kgDs7b7dx6XK0G1lPXFj6bEBG
TwitterAgent.sources.Twitter.consumerSecret=aoTZ0BkEZ3Xsm5ftAf6HGtr9qjzekLdQsQSWHdvZTHBnUzcPZQ
TwitterAgent.sources.Twitter.accessToken=3725631974-2B67v5IxJVQ0sLfKV80jyuSFm5G1qpcOBQp1Tba
TwitterAgent.sources.Twitter.accessTokenSecret=3IMjxnEBcpVZGLH7Ed4Hw87UAzIFDwnTbSdYBnBJ2KmDl
TwitterAgent.sources.Twitter.keywords=hadoop, bigdata, mapreduce, mahout, hbase, nosql
# Describing/Configuring the sink

TwitterAgent.sources.Twitter.keywords= hadoop,election,sports, cricket,Big data

TwitterAgent.sinks.HDFS.channel=MemChannel
TwitterAgent.sinks.HDFS.type=hdfs
TwitterAgent.sinks.HDFS.hdfs.path=hdfs://localhost:9000/user/flume/tweets
TwitterAgent.sinks.HDFS.hdfs.fileType=DataStream
TwitterAgent.sinks.HDFS.hdfs.writeformat=Text
TwitterAgent.sinks.HDFS.hdfs.batchSize=1000
TwitterAgent.sinks.HDFS.hdfs.rollSize=0
TwitterAgent.sinks.HDFS.hdfs.rollCount=10000
TwitterAgent.sinks.HDFS.hdfs.rollInterval=600

TwitterAgent.channels.MemChannel.type=memory
TwitterAgent.channels.MemChannel.capacity=10000
TwitterAgent.channels.MemChannel.transactionCapacity=1000

TwitterAgent.sources.Twitter.channels = MemChannel
TwitterAgent.sinks.HDFS.channel = MemChannel

```

### Step 8:

We have to decide which keywords tweet data to be collected from the twitter application. So, you can change the keywords in the TwitterAgent.source.Twitter.keywords command.

In our example, we are fetching tweet data related to stock market, recommendation, Hadoop, election, sports cricket and Big data.



```

TwitterAgent.sources = Twitter
TwitterAgent.channels = MemChannel
TwitterAgent.sinks = HDFS

# Describing/Configuring the source
TwitterAgent.sources.Twitter.type = org.apache.flume.source.twitter.TwitterSource
TwitterAgent.sources.Twitter.consumerKey=kgDs7b7dx6XK0G1lPXfj6bEBG
TwitterAgent.sources.Twitter.consumerSecret=aoTz0BkEZ3Xsm5ftAf6HGtr9qjzekLdQsQSWHdvZTHBnUzcPZQ
TwitterAgent.sources.Twitter.accessToken=3725631974-2B67v5IxJVQ0sLfKV80jyuSFm5GlpcoBQp1Tba
TwitterAgent.sources.Twitter.accessTokenSecret=3TlMlxnEBcV7G1H7Ed4Hw87UAzTFDwnTbSdYBnB12KmD1
TwitterAgent.sources.Twitter.keywords=stock, stock market, recommendation, hadoop, bigdata, mapreduce, mahout, hbase, nosql
# Describing/Configuring the sink

TwitterAgent.sources.Twitter.keywords= hadoop,election,sports, cricket,Big data

```

### Step 9:

Open a new terminal and start all the Hadoop daemons, before running the flume command to fetch the twitter data. Use the 'jps' command to see the running Hadoop daemons.

```

acadgild@localhost:~
File Edit View Search Terminal Help
[acadgild@localhost ~]$ start-all.sh
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
18/10/03 12:30:54 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /home/acadgild/install/hadoop/hadoop-2.
6.5/logs/hadoop-acadgild-namenode-localhost.localdomain.out
localhost: starting datanode, logging to /home/acadgild/install/hadoop/hadoop-2.
6.5/logs/hadoop-acadgild-datanode-localhost.localdomain.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/acadgild/install/hadoop/ha
doo-2.6.5/logs/hadoop-acadgild-secondarynamenode-localhost.localdomain.out
18/10/03 12:31:14 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/
logs/yarn-acadgild-resourcemanager-localhost.localdomain.out
localhost: starting nodemanager, logging to /home/acadgild/install/hadoop/hadoo
p-2.6.5/logs/yarn-acadgild-nodemanager-localhost.localdomain.out
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ jps
9463 NodeManager
9016 DataNode
9784 Jps
8888 NameNode
9210 SecondaryNameNode
9359 ResourceManager
[acadgild@localhost ~]$

```

### Step 10:

Now create a new directory inside HDFS path, where the Twitter tweet data should be stored.

```

acadgild@localhost:~
File Edit View Search Terminal Help
[acadgild@localhost ~]$ hadoop fs -mkdir -p /user/flume/tweets
18/10/03 12:37:15 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platf
orm... using builtin-java classes where applicable
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$ hadoop fs -ls /user/flume/
18/10/03 12:37:27 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platf
orm... using builtin-java classes where applicable
Found 1 items
drwxr-xr-x - acadgild supergroup          0 2018-10-03 12:37 /user/flume/tweets
[acadgild@localhost ~]$

```

### Step 11:

For fetching the data from Twitter, Use the below command to fetch the twitter tweet data into the HDFS cluster path.

**\$flume-ng agent -n TwitterAgent -f <location of created/edited conf file>**

```

acadgild@localhost:~
File Edit View Search Terminal Help
[acadgild@localhost ~]$ flume-ng agent -n TwitterAgent -f /home/acadgild/install/flume/apache-flu
me-1.8.0-bin/conf/flume.conf

```

### Step 12:

The above command will start fetching data from Twitter and streams it into the HDFS given path.



```
File Edit View Search Terminal Help
at org.apache.hadoop.io.retry.RetryInvocationHandler.invoke(RetryInvocationHandler.java:102)
at com.sun.proxy.$Proxy14.create(Unknown Source)
at org.apache.hadoop.hdfs.DFSOutputStream.newStreamForCreate(DFSOutputStream.java:1721)
at org.apache.hadoop.hdfs.DFSClient.create(DFSClient.java:1657)
at org.apache.hadoop.hdfs.DFSClient.create(DFSClient.java:1582)
at org.apache.hadoop.hdfs.DistributedFileSystem$6.doCall(DistributedFileSystem.java:397)
at org.apache.hadoop.hdfs.DistributedFileSystem$6.doCall(DistributedFileSystem.java:393)
at org.apache.hadoop.fs.FileSystemLinkResolver.resolve(FileSystemLinkResolver.java:81)
at org.apache.hadoop.hdfs.DistributedFileSystem.create(DistributedFileSystem.java:393)
at org.apache.hadoop.hdfs.DistributedFileSystem.create(DistributedFileSystem.java:337)
at org.apache.hadoop.fs.FileSystem.create(FileSystem.java:908)
at org.apache.hadoop.fs.FileSystem.create(FileSystem.java:889)
at org.apache.hadoop.fs.FileSystem.create(FileSystem.java:786)
at org.apache.hadoop.fs.FileSystem.create(FileSystem.java:775)
at org.apache.hadoop.hdfs.HDFSDataStream.doOpen(HDFSDataStream.java:81)
at org.apache.flume.sink.hdfs.HDFSDataStream.open(HDFSDataStream.java:108)
at org.apache.flume.sink.hdfs.BucketWriter$1.call(BucketWriter.java:262)
at org.apache.flume.sink.hdfs.BucketWriter$1.call(BucketWriter.java:252)
at org.apache.flume.sink.hdfs.BucketWriter$9$1.run(BucketWriter.java:701)
at org.apache.flume.auth.SimpleAuthenticator.execute(SimpleAuthenticator.java:50)
at org.apache.flume.sink.hdfs.BucketWriter$9.call(BucketWriter.java:698)
at java.util.concurrent.FutureTask.run(FutureTask.java:266)
at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1149)
at java.util.concurrent.ThreadPoolExecutor$Worker.run(ThreadPoolExecutor.java:624)
at java.lang.Thread.run(Thread.java:748)
Caused by: java.net.ConnectException: Connection refused
at sun.nio.ch.SocketChannelImpl.checkConnect(Native Method)
at sun.nio.ch.SocketChannelImpl.finishConnect(SocketChannelImpl.java:717)
at org.apache.hadoop.net.SocketIOWithTimeout.connect(SocketIOWithTimeout.java:206)
at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:530)
at org.apache.hadoop.net.NetUtils.connect(NetUtils.java:494)
at org.apache.hadoop.ipc.Client$Connection.setupConnection(Client.java:609)
at org.apache.hadoop.ipc.Client$Connection.setupIOstreams(Client.java:707)
at org.apache.hadoop.ipc.Client$Connection.access$2800(Client.java:370)
at org.apache.hadoop.ipc.Client.getConnection(Client.java:1523)
at org.apache.hadoop.ipc.Client.call(Client.java:1440)
... 33 more
```

### Step 13:

Once the tweet data started it into the given HDFS path we can use 'Ctrl+c' command to stop the streaming process.

### Step 14:

To check the contents of the tweet data we can use the following command:

***hadoop fs -ls /user/flume/tweets***



```
acadgild@localhost:~
File Edit View Search Terminal Tabs Help
acadgild@localhost:~
[acadgild@localhost ~]$ hadoop fs -ls /user/flume/tweets
18/10/03 18:17:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-
ava classes where applicable
Found 1 items
-rw-r--r-- 1 acadgild supergroup 8 2018-10-03 18:09 /user/flume/tweets/FlumeData.1538569010037
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$
```

We can see that the all the streaming tweets are get stored in the FlumeData.xxx file in hadoop.